

## 1 Problem 1: Getting Started with the RAMA-2200

In this homework, you will be using the RAMA-2200 ISA to complete a GCD function. Before you begin, you should familiarize yourself with the available instructions, the register conventions and the calling convention of RAMA-2200. Details can be found in the section, Appendix A: RAMA-2200 Instruction Set Architecture, at the end of this document.

The `assembly` folder contains several tools for you to use:

- `assembler.py`: a basic assembler that can take your assembly code and convert it into binary instructions for the RAMA-2200.
- `rama2200.py`: the ISA definition file for the assembler, which tells `assembler.py` the instructions supported by the RAMA-2200 and their formats.
- `rama2200-sim.py`: A simulator of the RAMA-2200 machine. The simulator reads binary instructions and emulates the RAMA-2200 machine, letting you single-step through instructions and check their results.

To learn how to run these tools, see the `README.md` file in the `assembly` directory.

Before you begin work on the second problem of the homework, try writing a simple program for the RAMA-2200 architecture. This should help you familiarize yourself with the available instructions.

We have provided a template, `mod.s`, for you to use for this purpose. Try writing a program that performs the `mod` operation on the two provided arguments. A correct implementation will result in a value of 2.

You can use the following C code snippet as a guide to implement this function:

```
int mod(int a, int b) {
    int x = a;
    while (x >= b) {
        x = x - b;
    }
    return x;
}
```

There is no turn-in for this portion of the assignment, but it is **recommended** that you attempt it in order to familiarize yourself with the ISA.

## 2 Problem 2: GCD Program

For this problem, you will be implementing the missing portions of a GCD algorithm we have provided for you.

You'll be finishing a **recursive** implementation of GCD that follows the RAMA-2200 calling convention. Recursive functions always obtain a return address through the function call and return to the callee using the return address.

**You must use the stack pointer (\$sp) and frame pointer (\$fp) registers as described in the textbook and lecture slides.**

Here is the C code for the GCD algorithm you have been provided:

```
int gcd(int a, int b) {
    if (b == 0) {
        return a;
    }
    return gcd(b, a % b);
}
```

Open the `gcd.s` file in the `assembly` directory. This file contains an implementation of the GCD algorithm that is missing significant portions of the calling convention.

1. Complete the code below the `main` label to initialize the stack pointer (\$sp) to point at `0xFFFF`. You should use the provided “stack” to load this value into the register.
2. Complete the given GCD subroutine by implementing the various missing portions of the RAMA-2200 calling convention. Each location where you need to implement a portion of the calling convention is marked with a `TODO` label as well as a short hint describing the portion of the calling convention you should be implementing.

Please note that we will be testing your implementation for multiple values of  $a$  and  $b$ , so please do not attempt to hardcode your solutions. We will only be testing for the cases where  $a \geq b$ .

## 3 Problem 3: Short Answer

Please answer the following question in the file named `gcd.s`:

1. The RAMA-2200 instruction set contains an instruction called `jalr` that is used to jump to a location while saving a return address. However, this functionality could be emulated using a combination of other instructions available in the ISA. Describe a sequence of other instructions in the RAMA-2200 ISA that you may use to accomplish the functionality of `jalr`.

## 4 Deliverables

- `gcd.s`: your GCD assembly code from Problem 2: GCD Program
- `answers.txt`: your answer to Problem 3: Short Answer

Submit these files to Canvas before the assignment deadline.

The TAs should be able to type `python assembler.py -i rama2200 --sym gcd.s` and then `python rama2200-sim.py gcd.bin` to run your code. If you cannot do this with your submission, then you have done something wrong.

## 5 Appendix A: RAMA-2200 Instruction Set Architecture

The RAMA-2200 is a simple, yet capable computer architecture. The RAMA-2200 combines attributes of both ARM and the LC-2200 ISA defined in the Ramachandran & Leahy textbook for CS 2200.

The RAMA-2200 is a **word-addressable, 32-bit** computer. **All addresses refer to words**, i.e. the first word (four bytes) in memory occupies address 0x0, the second word, 0x1, etc.

All memory addresses are truncated to 16 bits on access, discarding the 16 most significant bits if the address was stored in a 32-bit register. This provides roughly 64 KB of addressable memory.

### 5.1 Registers

The RAMA-2200 has 16 general-purpose registers. While there are no hardware-enforced restraints on the uses of these registers, your code is expected to follow the conventions outlined below.

Table 1: Registers and their Uses

Register Number	Name	Use	Callee Save?
0	\$zero	Always Zero	NA
1	\$at	Assembler/Target Address	NA
2	\$v0	Return Value	No
3	\$a0	Argument 1	No
4	\$a1	Argument 2	No
5	\$a2	Argument 3	No
6	\$t0	Temporary Variable	No
7	\$t1	Temporary Variable	No
8	\$t2	Temporary Variable	No
9	\$s0	Saved Register	Yes
10	\$s1	Saved Register	Yes
11	\$s2	Saved Register	Yes
12	\$k0	Reserved for OS and Traps	NA
13	\$sp	Stack Pointer	No
14	\$fp	Frame Pointer	Yes
15	\$ra	Return Address	No

1. **Register 0** is always read as zero. Any values written to it are discarded. **Note:** for the purposes of this project, you must implement the zero register. Regardless of what is written to this register, it should always output zero.
2. **Register 1** is used to hold the target address of a jump. It may also be used by pseudo-instructions generated by the assembler.
3. **Register 2** is where you should store any returned value from a subroutine call.
4. **Registers 3 - 5** are used to store function/subroutine arguments. **Note:** registers 2 through 8 should be placed on the stack if the caller wants to retain those values. These registers are fair game for the callee (subroutine) to trash.
5. **Registers 6 - 8** are designated for temporary variables. The caller must save these registers if they want these values to be retained.
6. **Registers 9 - 11** are saved registers. The caller may assume that these registers are never tampered with by the subroutine. If the subroutine needs these registers, then it should place them on the stack and restore them before they jump back to the caller.

7. **Register 12** is reserved for handling interrupts. While it should be implemented, it otherwise will not have any special use on this assignment.
8. **Register 13** is your anchor on the stack. It keeps track of the top of the activation record for a subroutine.
9. **Register 14** is used to point to the first address on the activation record for the currently executing process.
10. **Register 15** is used to store the address a subroutine should return to when it is finished executing.

## 5.2 Instruction Overview

The RAMA-2200 supports a variety of instruction forms, only a few of which we will use for this project. The instructions we will implement in this project are summarized below.

Table 2: RAMA-2200 Instruction Set

	31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
ADD	0000								DR																								SR2
NAND	0001								DR																								SR2
ADDI	0010								DR																								immval20
LW	0011								DR																								offset20
SW	0100								SR																								offset20
BEQ	0101								SR1																								offset20
JALR	0110								RA																								unused
HALT	0111																																unused
BLT	1000								SR1																								offset20
LEA	1001								DR																								offset20

### 5.2.1 Conditional Branching

Conditional branching in the RAMA-2200 ISA is provided via the BEQ (“branch if equal”) and BLT (“branch if less than”) instructions. BEQ will branch to address “incrementedPC + offset20” only if SR1 and SR2 are equal. Meanwhile, BLT will branch to address “incrementedPC + offset20” only if SR1 is less than SR2.

## 5.3 Detailed Instruction Reference

### 5.3.1 ADD

#### Assembler Syntax

ADD DR, SR1, SR2

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0000	DR	SR1	unused																												SR2

#### Operation

DR = SR1 + SR2;

#### Description

The ADD instruction obtains the first source operand from the SR1 register. The second source operand is obtained from the SR2 register. The second operand is added to the first source operand, and the result is stored in DR.

### 5.3.2 NAND

#### Assembler Syntax

NAND DR, SR1, SR2

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0001	DR	SR1	unused																												SR2

#### Operation

DR = ~(SR1 & SR2);

#### Description

The NAND instruction performs a logical NAND (AND NOT) on the source operands obtained from SR1 and SR2. The result is stored in DR.

**HINT:** A logical NOT can be achieved by performing a NAND with both source operands the same. For instance,

NAND DR, SR1, SR1

...achieves the following logical operation:  $DR \leftarrow \overline{SR1}$ .

### 5.3.3 ADDI

#### Assembler Syntax

ADDI DR, SR1, immval20

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0010				DR				SR1				immval20																			

#### Operation

DR = SR1 + SEXT(immval20);

#### Description

The ADDI instruction obtains the first source operand from the SR1 register. The second source operand is obtained by sign-extending the immval20 field to 32 bits. The resulting operand is added to the first source operand, and the result is stored in DR.

### 5.3.4 LW

#### Assembler Syntax

LW DR, offset20(BaseR)

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0011				DR				BaseR				offset20																			

#### Operation

DR = MEM[BaseR + SEXT(offset20)];

#### Description

An address is computed by sign-extending bits [19:0] to 32 bits and then adding this result to the contents of the register specified by bits [23:20]. The 32-bit word at this address is loaded into DR.

### 5.3.5 SW

#### Assembler Syntax

SW SR, offset20(BaseR)

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0100				SR				BaseR				offset20																			

#### Operation

MEM[BaseR + SEXT(offset20)] = SR;

#### Description

An address is computed by sign-extending bits [19:0] to 32 bits and then adding this result to the contents of the register specified by bits [23:20]. The 32-bit word obtained from register SR is then stored at this address.

### 5.3.6 BEQ

#### Assembler Syntax

BEQ SR1, SR2, offset20

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0101				SR1				SR2				offset20																			

#### Operation

```
if (SR1 == SR2) {
    PC = incrementedPC + offset20
}
```

#### Description

A branch is taken if SR1 and SR2 are equal. If this is the case, the PC will be set to the sum of the incremented PC (since we have already undergone fetch) and the sign-extended offset[19:0].

### 5.3.7 JALR

#### Assembler Syntax

JALR RA, AT

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0110				RA				AT				unused																			

#### Operation

RA = PC;

PC = AT;

#### Description

First, the incremented PC (address of the instruction + 1) is stored into register RA. Next, the PC is loaded with the value of register AT, and the computer resumes execution at the new PC.

### 5.3.8 HALT

#### Assembler Syntax

HALT

#### Encoding

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0111				unused																											

#### Description

The machine is brought to a halt and executes no further instructions.



**5.3.9 BLT****Assembler Syntax**

BLT      SR1, SR2, offset20

**Encoding**

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1000				SR1				SR2				offset20																			

**Operation**

```
if (SR1 < SR2) {
    PC = incrementedPC + offset20
}
```

**Description**

A branch is taken if SR1 is less than SR2. If this is the case, the PC will be set to the sum of the incremented PC (since we have already undergone fetch) and the sign-extended offset[19:0].

**5.3.10 LEA****Assembler Syntax**

LEA      DR, label

**Encoding**

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1001				DR				unused				PCoffset20																			

**Operation**

DR = PC + SEXT(PCoffset20);

**Description**

An address is computed by sign-extending bits [19:0] to 32 bits and adding this result to the incremented PC (address of instruction + 1). It then stores the computed address into register DR.